

In [1]:

```
import pandas as pd

import numpy as np

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.linear_model import Ridge, RidgeCV, Lasso

from sklearn.preprocessing import StandardScaler
```

In [4]:

```
#data

data=pd.read_csv(r"C:\Users\mural\Downloads\fiat500_VehicleSelection_Dataset.csv")

data
```

Out[4]:

|      | ID   | model  | engine_power | age_in_days | km     | previous_owners | lat       |        |
|------|------|--------|--------------|-------------|--------|-----------------|-----------|--------|
| 0    | 1    | lounge | 51           | 882         | 25000  | 1               | 44.907242 | 8.611  |
| 1    | 2    | pop    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241 |
| 2    | 3    | sport  | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417 |
| 3    | 4    | lounge | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634 |
| 4    | 5    | pop    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495 |
| ...  | ...  | ...    | ...          | ...         | ...    | ...             | ...       | ...    |
| 1533 | 1534 | sport  | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704  |
| 1534 | 1535 | lounge | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666  |
| 1535 | 1536 | pop    | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413  |
| 1536 | 1537 | lounge | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682  |
| 1537 | 1538 | pop    | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568 |

1538 rows × 9 columns



In [5]:

```
data.head()
```

Out[5]:

|   | ID | model  | engine_power | age_in_days | km     | previous_owners | lat       | lon       |
|---|----|--------|--------------|-------------|--------|-----------------|-----------|-----------|
| 0 | 1  | lounge | 51           | 882         | 25000  | 1               | 44.907242 | 8.611560  |
| 1 | 2  | pop    | 51           | 1186        | 32500  | 1               | 45.666359 | 12.241890 |
| 2 | 3  | sport  | 74           | 4658        | 142228 | 1               | 45.503300 | 11.417840 |
| 3 | 4  | lounge | 51           | 2739        | 160000 | 1               | 40.633171 | 17.634609 |
| 4 | 5  | pop    | 73           | 3074        | 106880 | 1               | 41.903221 | 12.495650 |

In [6]:

```
data.tail()
```

Out[6]:

|      | ID   | model  | engine_power | age_in_days | km     | previous_owners | lat       | lon    |
|------|------|--------|--------------|-------------|--------|-----------------|-----------|--------|
| 1533 | 1534 | sport  | 51           | 3712        | 115280 | 1               | 45.069679 | 7.704  |
| 1534 | 1535 | lounge | 74           | 3835        | 112000 | 1               | 45.845692 | 8.666  |
| 1535 | 1536 | pop    | 51           | 2223        | 60457  | 1               | 45.481541 | 9.413  |
| 1536 | 1537 | lounge | 51           | 2557        | 80750  | 1               | 45.000702 | 7.682  |
| 1537 | 1538 | pop    | 51           | 1766        | 54276  | 1               | 40.323410 | 17.568 |

In [12]:


```
data.drop(columns=["model"],inplace=True)
```

In [13]:

```
plt.figure(figsize=(10,10))  
sns.heatmap(data.corr(),annot = True)
```

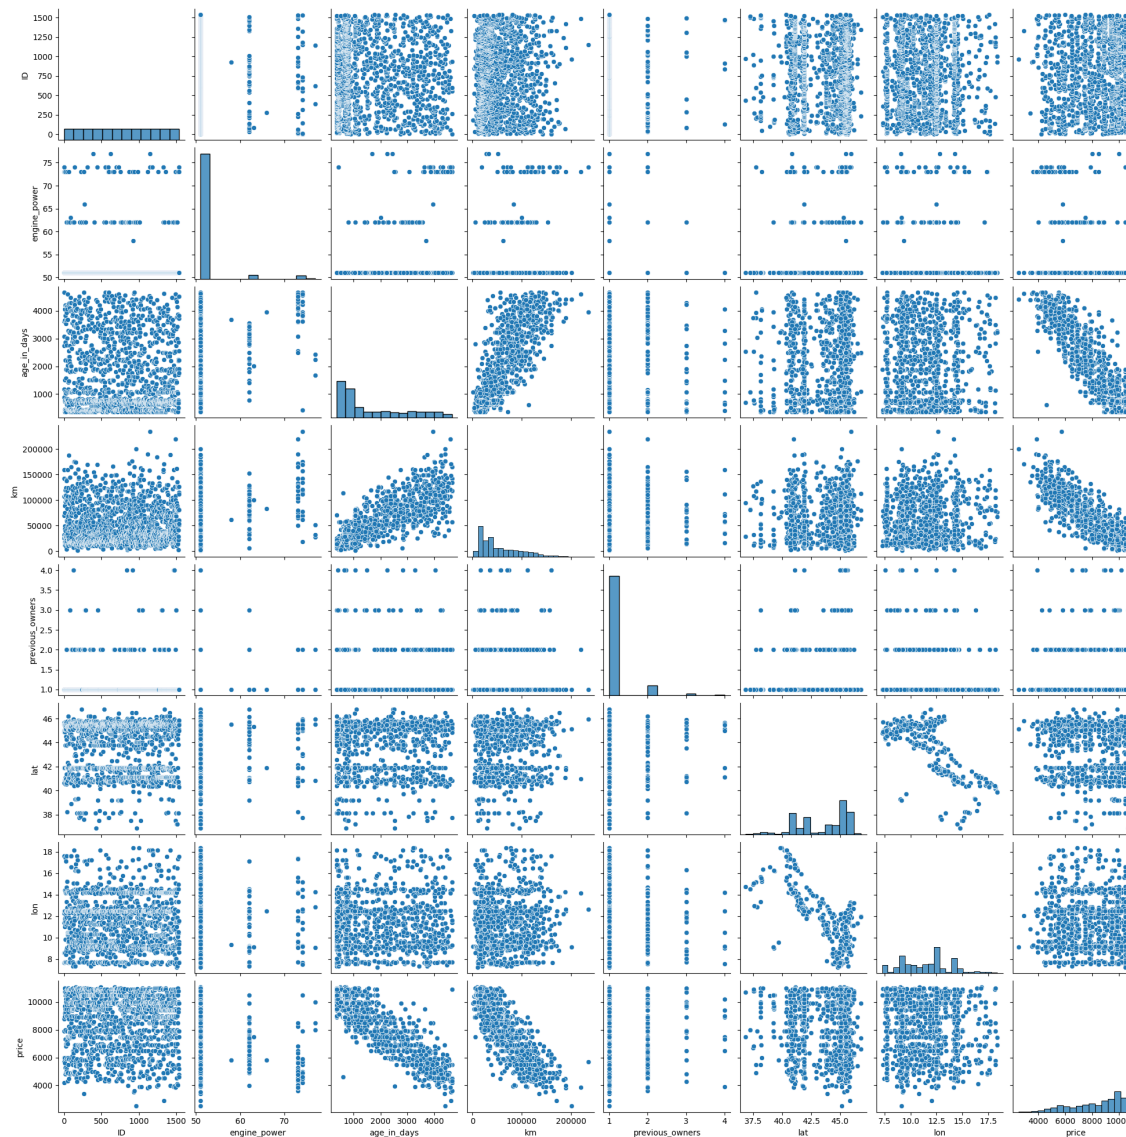
Out[13]:

&lt;Axes: &gt;



In [15]:

```
sns.pairplot(data)
data.price = np.log(data.price)
```



In [16]:

```
features = data.columns[0:2]
target = data.columns[-1]
#X and y values
X = data[features].values
y = data[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (1076, 2)  
 The dimension of X\_test is (462, 2)

In [17]:

```
#Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07906758951709636

The test score for lr model is 0.08573839649638304

In [19]:

```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```


Ridge Model:

The train score for ridge model is 0.07906120163510788

The test score for ridge model is 0.08541192691344546

In [20]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```




In [22]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[22]:

<Axes: >



In [23]:

```
#Using the linear CV model

from sklearn.linear_model import LassoCV

#Lasso Cross validation

lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)

#score

print(lasso_cv.score(X_train, y_train))

print(lasso_cv.score(X_test, y_test))
```

```
0.07906730311134957
0.08575009503364805
```

## ELASTICNET REGRESSION

In [24]:

```
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[ 1.60035419e-05 -0.00000000e+00]
9.013884247807239
```

In [25]:

```
y_pred_elastic=regr.predict(X_train)
```

In [26]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 0.06540870300958243
```

In [ ]: